

# HOW TO OPTIMIZE

# How to Optimize

## Overview

- ▶ Programmer's guide
  - ▶ General coding advice
  - ▶ C++ coding pitfalls
  - ▶ Target specific issues
- ▶ How to (profile on) Lauterbach (target debuggers)
  - ▶ Setup
  - ▶ Interpreting results
- ▶ Advanced topics
  - ▶ Optimization strategy
  - ▶ Local: reading the assembly code
  - ▶ Global: processor objects

*Don't optimize just yet,  
clean code is usually fast.*

# Programmer's Guide

## General coding advice (1/2)

- ▶ Golden Rule: Code gets faster if it has less to do
  - ▶ Your code does lots of repeated work → get rid of that
  - ▶ Sometimes, a better algorithm using more of the available information is necessary.
- ▶ Factor out common code
  - ▶ In loop headers: `size()`, `end()` etc. are often needlessly repeated
  - ▶ Any (sub-)expression that does not depend on inner loop state should be moved out
  - ▶ Common sub-expressions can also be found in sequential code, particularly function results

```
x = foo()[0]; y = foo()[1];           // foo constructs an expensive object twice
v = foo(); x = vec[0]; y = vec[1];    // only 1 expensive operation
```

# Programmer's Guide

## General coding advice (2/2)

- ▶ “Big-O” runtime complexity: beware of nested loops
  - ▶ Number of operations is the product of all iteration counts, e.g.  $O(\text{numLocations} * \text{numObjects} * \text{numNewLocations})$
  - ▶ Cutting the iterations short by skipping iterations with `if()` often less effective in “dynamic situations”  
→ extreme peaks while average runtimes look good
  - ▶ The inner loop(s) may be “hidden” in sub-functions  
Some container classes do a full scan for e.g. `a.get_constObjectRef(id)`.
  - ▶ Most can be fixed by sorting before matching
- ▶ Order of checks in hot loops is important
  - ▶ Put the most restrictive check first
  - ▶ Make sure that code gets actually inlined (in case the check is in some sub-function)
  - ▶ Simplify the hot loop + main check combination as much as possible (ideal: pointer walk + comparison)

# Programmer's Guide

## C++ coding pitfalls (1/2)

- ▶ The following are rules, make exceptions at your discretion
- ▶ Function parameters
  - ▶ Pass non-trivial parameters by const reference (basically everything that is not a single number or pointer)
  - ▶ Pass trivial parameters by value
  - ▶ Avoid default values for non-trivial parameters
- ▶ **Small** virtual functions will usually not be inlined → see that you can drop the “virtual” for those

# Programmer's Guide

## C++ coding pitfalls (2/2)

### ► Function results

- Getters for internal state should return non-trivial data by const reference.

- Avoid intermediate copies

```
x = SomeType(MyFunction()); // usually, a temporary object will be created
x = MyFunction();           // no additional object, also robust against signature changes
```

- Take advantage of the RVO (Return Value Optimization)

Kicks in, if the receiving / calling side is constructing a new object

```
SomeType x; // construct x
x = MyFunction(); // creates a temp object, copies it to x, temp object is destroyed
SomeType y = MyFunction(); // RVO: function constructs its result directly in y
```

# Programmer's Guide

## Target specific issues

### ► Float vs. double

- Target has 32 bit float support in hardware (FPU)
- 64 bit floats are software-only
- Be sure to use `float`, `float32_t` instead of `double`, `float64_t`
- `0.5` is a double, `0.5f` is a float

### ► No inlining of loops

- Any function containing a `for()` loop will not be inlined, even if the optimizer can completely eliminate the loop as it often happens in templates.
- Maybe use recursion instead but only for loops that are **short and** have a **fixed** number of iterations

### ► Optimizer will not factor out repeated calculations, especially floating point functions, e.g. `pow` in

```
for(i=0;i<4;++i) {x[i] += pow(globalVariable);}
```



*Wer misst, misst Mist.  
(He who measures, measures rubbish.)*

# How to Lauterbach Setup

- ▶ Make performance tests repeatable: automate input and output
  - ▶ Restbus simulation may be sufficient for testing COM stack performance
  - ▶ Sensor data might be feed in through HiL but that is complicated
  - ▶ Bypass as alternative: extract data from MDF, store it in flash, replay by filling Daddy ports  
(see copy runnables in `jobCore1_T_BG()` in `per_profiling` branch in `pj_vwmqb37w.git`)
- ▶ Make outputs verifiable, e.g. the following way:
  - ▶ At runtime, store key info like DEP & SEP object counts in an array; one entry per task cycle
  - ▶ Dump the data to file, import in Excel and quick check for deviations from previous runs
  - ▶ Lauterbach script to dump the `scom::perfLog` array:  

```
PRinTer.FILE C:\sbx\PerfLog.csv CSV  
WinPrint.var.fixedtable %LOCATION.off %hex.off scom::perfLog
```

# How to Lauterbach

## Interpreting results

► Use `B::perf.ListFuncMod /core 1`

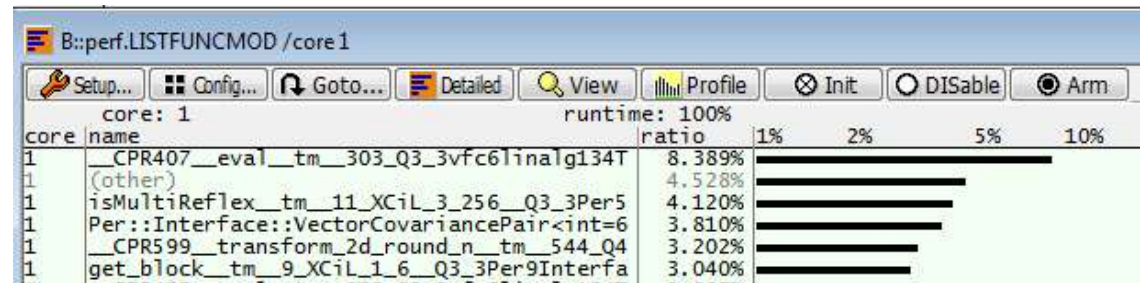
► All perf views show the same data

► Use Arm / Disable buttons to start / stop

► Often shows mangled C names

► “(other)” are system functions like `memcpy`, `sqrt`, `__ll_div` for 64 bit float math

► `B::perf.List /core 1` will show those in detail but as “code ranges”



1	Per::Interface::VectorCovariancePair<int=8>::set_matrix--get_subVector	2.095%	=====
1	sqrtf--sinf	2.034%	=====
1	get_block__tm_9_XCiL_1_6_Q3_3Per9Interface36VectorCovariancePair__tm	1.717%	=====

► Results do not cover costs of sub-function execution

► You may stop the execution at any point, set breakpoints etc.

► Double click on perf list entry will bring the respective function up

► Use “mixed” view to be pointed to the correct template instance

*If brute force is not the answer,  
you are just not using enough of it.*

# Advanced Topics

## Optimization strategy

### ► **Measure before you try to optimize!**

- Successful optimization is an art form like good management

### ► Develop a Theory of Computation for the respective SW / HW combination

- A naïve, simplistic mental model of what the system has to do, e.g.  
PER = read locations, group to objects, extrapolate movement, match with new locations
  - What would the minimal amount of operations be, that is needed to produce the desired result?
  - What is “core functionality” and what is “overhead”?
- ### ► Generic strategy: Finding and understanding issues is harder than fixing them
- Gain system understanding: Identify frequently used low-level functions and “squeeze” them hard
  - Identify “big-O” mistakes: Use call counts (callgrind!) to identify high “fan-outs” to sub-functions
  - Reduce friction, improve SNR: Either no sub-functions after inlining or runtime dominated by sub-functions

# Advanced Topics

## Local: reading the assembly code (1/2)

- ▶ The Lauterbach debugger will show you the machine instructions
  - ▶ Often, you can identify hot spots and inefficiencies directly
  - ▶ Use the compiler-generated assembly listing to see whether new code is better
- ▶ Search the `.lst` files for the implementation of the code lines in questions
- ▶ Things check at assembly level:
  - ▶ Has a particular function been inlined or even completely eliminated?
  - ▶ Are there a virtual function calls on the hot path, e.g. `call [a2]`?
  - ▶ Is there repetitive / superfluous code (see next slide)?
  - ▶ What would be the minimal set of assembly operations for the hot path?  
Try to express those as C/C++ code.

# Advanced Topics

## Local: reading the assembly code (2/2)

- ▶ GHS has a tendency to repeat address calculations; use iterators in that case
- ▶ Example: one iteration in an unrolled loop of matrix operations

```
l_sum -= f_L(i,k) * f_L(j,k);
```

```
msub.f    d4,d4,d0,d6
ld16.w    d0,[a12]
addsc.a    a12,a13,d1,0
lea        a12,[a12]0x14
ld16.w    d6,[a12]
addsc.a    a12,a13,d5,0
lea        a12,[a12]0x10
msub.f    d4,d4,d0,d6
ld16.w    d0,[a12]
```

*actually part of  
the next iteration  
but the pattern  
still holds*

```
l_sum -= *l_ik_it * *l_jk_it;
```

```
msub.f    d4,d4,d0,d1
ld.w       d1,[a3]-20
ld.w       d0,[a2]-20
msub.f    d4,d4,d0,d1
ld.w       d1,[a3]-24
```

# Advanced Topics

## Global: processing objects

- ▶ Extension of the “factor out common code” tactics for this scenario
  - ▶ One or more deeply nested functions are repeatedly called by some function high up in the stack
  - ▶ Nested code does not “know” about earlier executions
  - ▶ Nested code needs to re-compute things over and over
- ▶ Solution: Add a parameter that preserves the context between calls
  - ▶ Could be a caching object to short-circuit lookups
  - ▶ Could be data pre-calculated in the higher up function
- ▶ Enhanced version: Processing objects
  - ▶ Turn the context object into something semantically meaningful, e.g. `TMyTransformation`
  - ▶ Move the logic of the nested function into methods of the new object
  - ▶ Optimize code further between those methods; they are probably tightly linked